

In [20]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
```

In [21]:

```
iris = pd.read_csv('C:\datasets\iris.csv')
```

In [22]:

```
iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Id              150 non-null   int64
 1   SepalLengthCm   150 non-null   float64
 2   SepalWidthCm    150 non-null   float64
 3   PetalLengthCm   150 non-null   float64
 4   PetalWidthCm    150 non-null   float64
 5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

In [23]:

```
iris.head()
```

Out[23]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

In [24]:

```
iris.tail()
```

Out[24]:

	<b>Id</b>	<b>SepalLengthCm</b>	<b>SepalWidthCm</b>	<b>PetalLengthCm</b>	<b>PetalWidthCm</b>	<b>Species</b>
<b>145</b>	146	6.7	3.0	5.2	2.3	Iris-virginica
<b>146</b>	147	6.3	2.5	5.0	1.9	Iris-virginica
<b>147</b>	148	6.5	3.0	5.2	2.0	Iris-virginica
<b>148</b>	149	6.2	3.4	5.4	2.3	Iris-virginica
<b>149</b>	150	5.9	3.0	5.1	1.8	Iris-virginica

In [25]:

```
iris.count()
```

Out[25]:

```
Id                150
SepalLengthCm    150
SepalWidthCm     150
PetalLengthCm    150
PetalWidthCm     150
Species          150
dtype: int64
```

In [26]:

```
iris.shape
```

Out[26]:

```
(150, 6)
```

In [27]:

```
iris.columns
```

Out[27]:

```
Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
      'Species'],
      dtype='object')
```

In [28]:

```
iris.describe()
```

Out[28]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
<b>count</b>	150.000000	150.000000	150.000000	150.000000	150.000000
<b>mean</b>	75.500000	5.843333	3.054000	3.758667	1.198667
<b>std</b>	43.445368	0.828066	0.433594	1.764420	0.763161
<b>min</b>	1.000000	4.300000	2.000000	1.000000	0.100000
<b>25%</b>	38.250000	5.100000	2.800000	1.600000	0.300000
<b>50%</b>	75.500000	5.800000	3.000000	4.350000	1.300000
<b>75%</b>	112.750000	6.400000	3.300000	5.100000	1.800000
<b>max</b>	150.000000	7.900000	4.400000	6.900000	2.500000

In [29]:

```
iris.groupby('Species').size()
```

Out[29]:

```
Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64
```

In [30]:

```
iris.sample(10)
```

Out[30]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
<b>46</b>	47	5.1	3.8	1.6	0.2	Iris-setosa
<b>14</b>	15	5.8	4.0	1.2	0.2	Iris-setosa
<b>119</b>	120	6.0	2.2	5.0	1.5	Iris-virginica
<b>136</b>	137	6.3	3.4	5.6	2.4	Iris-virginica
<b>23</b>	24	5.1	3.3	1.7	0.5	Iris-setosa
<b>34</b>	35	4.9	3.1	1.5	0.1	Iris-setosa
<b>100</b>	101	6.3	3.3	6.0	2.5	Iris-virginica
<b>84</b>	85	5.4	3.0	4.5	1.5	Iris-versicolor
<b>1</b>	2	4.9	3.0	1.4	0.2	Iris-setosa
<b>103</b>	104	6.3	2.9	5.6	1.8	Iris-virginica

# Removing the un-necessary column from the dataset

In [31]:

```
iris.drop("Id", axis=1, inplace = True)
```

In [32]:

```
iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   SepalLengthCm   150 non-null    float64
1   SepalWidthCm    150 non-null    float64
2   PetalLengthCm   150 non-null    float64
3   PetalWidthCm    150 non-null    float64
4   Species         150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [33]:

```
iris.head()
```

Out[33]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [34]:

```
iris.isnull()
```

Out[34]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...	...	...	...	...	...
145	False	False	False	False	False
146	False	False	False	False	False
147	False	False	False	False	False
148	False	False	False	False	False
149	False	False	False	False	False

150 rows × 5 columns


# EDA Analysis

In [35]:

```
fig = iris[iris.Species == 'Iris-setosa'].plot(kind='scatter', x='SepalLengthCm', y='SepalWidthCm')
iris[iris.Species == 'Iris-versicolor'].plot(kind='scatter', x='SepalLengthCm', y='SepalWidthCm')
iris[iris.Species == 'Iris-virginica'].plot(kind='scatter', x='SepalLengthCm', y='SepalWidthCm')

fig.set_xlabel('Sepal Length')
fig.set_ylabel('Sepal Width')
fig.set_title('Sepal Length Vs Width')

fig=plt.gcf()
fig.set_size_inches(10, 7)
plt.show()
```



In [36]:


```
sns.FacetGrid(iris, hue='Species', size=5)\
    .map(plt.scatter, 'SepalLengthCm', 'SepalWidthCm')\
    .add_legend()
```

C:\ProgramData\Anaconda\lib\site-packages\seaborn\axisgrid.py:316: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

```
warnings.warn(msg, UserWarning)
```

Out[36]:

<seaborn.axisgrid.FacetGrid at 0x29995f39700>




In [37]:

```
fig = iris[iris.Species == 'Iris-setosa'].plot(kind='scatter', x='PetalLengthCm', y='PetalWidthCm')
iris[iris.Species == 'Iris-versicolor'].plot(kind='scatter', x='PetalLengthCm', y='PetalWidthCm')
iris[iris.Species == 'Iris-virginica'].plot(kind='scatter', x='PetalLengthCm', y='PetalWidthCm')

fig.set_xlabel('Petal Length')
fig.set_ylabel('Petal Width')
fig.set_title('Petal Length Vs Width')

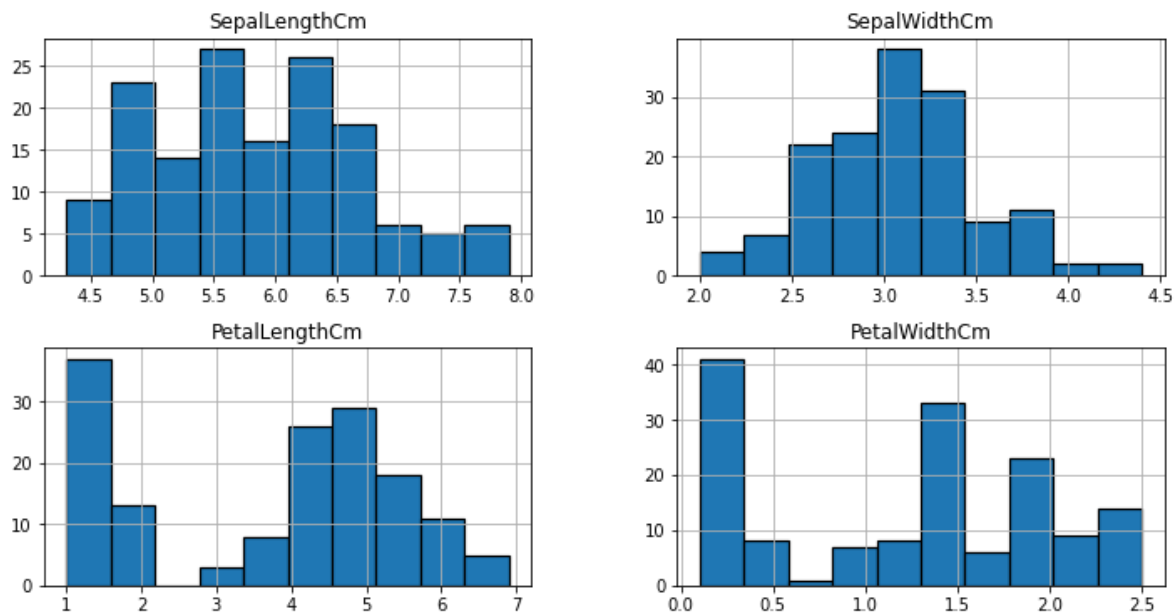
fig=plt.gcf()
fig.set_size_inches(10, 7)
plt.show()
```





In [38]:

```
iris.hist(edgecolor='black', linewidth=1.2)
fig = plt.gcf()
fig.set_size_inches(12,6)
plt.show()
```



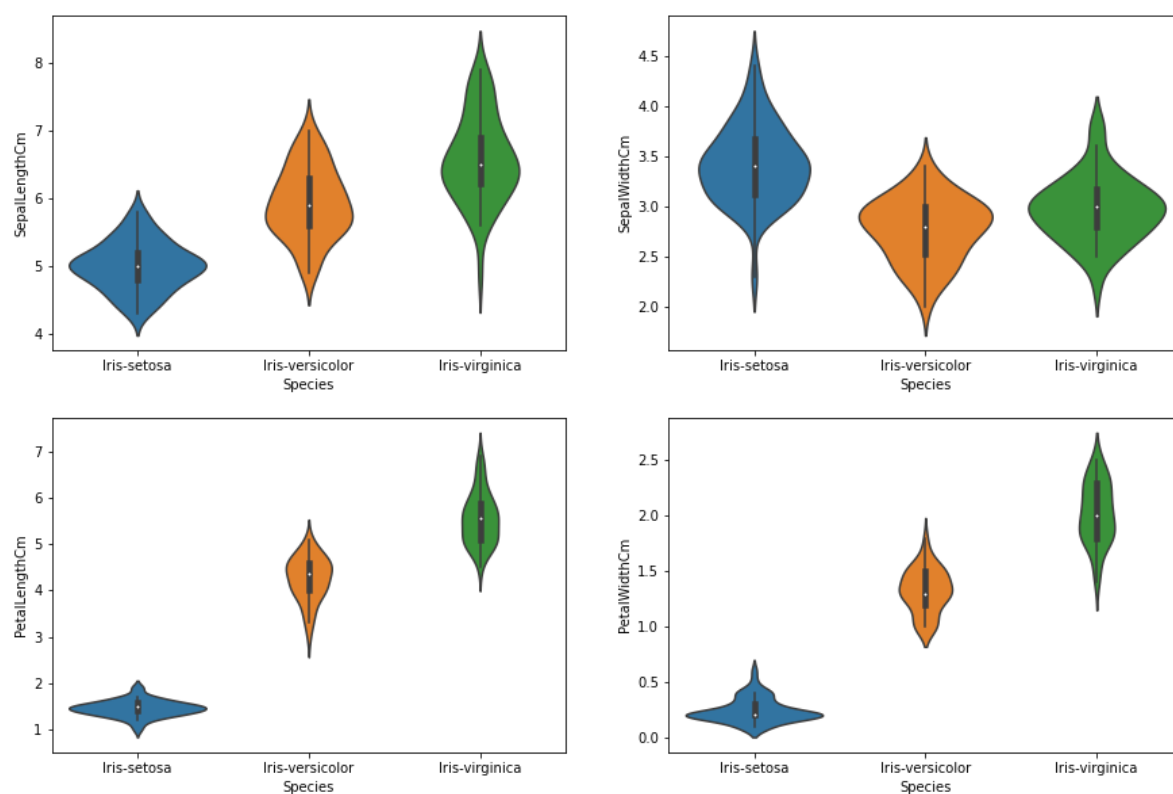
In [39]:

```
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.violinplot(x='Species', y = 'SepalLengthCm', data=iris)
plt.subplot(2,2,2)
sns.violinplot(x='Species', y = 'SepalWidthCm', data=iris)

plt.subplot(2,2,3)
sns.violinplot(x='Species', y = 'PetalLengthCm', data=iris)
plt.subplot(2,2,4)
sns.violinplot(x='Species', y = 'PetalWidthCm', data=iris)
```

Out[39]:

<AxesSubplot:xlabel='Species', ylabel='PetalWidthCm'>



In [40]:

```

from sklearn import svm # for SVM Algorithm
from sklearn.linear_model import LogisticRegression # for Logistic Regression Algorithm
from sklearn.model_selection import train_test_split # to split the dataset for training and testing
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

```

In [41]:

```

train, test = train_test_split(iris, test_size=0.3) # our main data split into train and test
# the attribute test_size=0.3 splits the data into 70% and 30% ratio. train=70% and test=30%
print(train.shape)
print(test.shape)

```

```

(105, 5)
(45, 5)

```

In [42]:

```

train_X = train[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']] # taking features from train data
train_y = train.Species # output of the training data

test_X = test[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']] # taking features from test data
test_y = test.Species # output value of the test data

```

In [43]:

```
train_X.head()
```

Out[43]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
74	6.4	2.9	4.3	1.3
28	5.2	3.4	1.4	0.2
147	6.5	3.0	5.2	2.0
94	5.6	2.7	4.2	1.3
34	4.9	3.1	1.5	0.1

In [44]:

```
test_X.head()
```

Out[44]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
125	7.2	3.2	6.0	1.8
29	4.7	3.2	1.6	0.2
7	5.0	3.4	1.5	0.2
20	5.4	3.4	1.7	0.2
57	4.9	2.4	3.3	1.0

In [45]:

```
train_y.head()
```

Out[45]:

```
74    Iris-versicolor
28      Iris-setosa
147    Iris-virginica
94    Iris-versicolor
34      Iris-setosa
Name: Species, dtype: object
```

## Decision Tree

In [48]:

```

model = DecisionTreeClassifier()
model.fit(train_X, train_y)
prediction = model.predict(test_X)
print('The accuracy of Decision Tree is: ', metrics.accuracy_score(prediction, test_y))
print('The confusion matrix for decision tree is: ', metrics.confusion_matrix(prediction, test_y))
print('The classification matrix for decision tree is: ', metrics.classification_report(test_y, prediction))

```

The accuracy of Decision Tree is: 0.9777777777777777

The confusion matrix for decision tree is: [[19 0 0]

[ 0 14 0]

[ 0 1 11]]

The classification matrix for decision tree is:

precision

recall f1-score support

Iris-setosa	1.00	1.00	1.00	19
Iris-versicolor	1.00	0.93	0.97	15
Iris-virginica	0.92	1.00	0.96	11
accuracy			0.98	45
macro avg	0.97	0.98	0.97	45
weighted avg	0.98	0.98	0.98	45

## Logistic Regression

In [47]:

```

model = LogisticRegression()
model.fit(train_X, train_y)
prediction = model.predict(test_X)
print('The accuracy of Decision Tree is: ', metrics.accuracy_score(prediction, test_y))
print('The confusion matrix for decision tree is: ', metrics.confusion_matrix(prediction, test_y))
print('The classification matrix for decision tree is: ', metrics.classification_report(test_y, prediction))

```

The accuracy of Decision Tree is: 0.9777777777777777

The confusion matrix for decision tree is: [[19 0 0]

[ 0 14 0]

[ 0 1 11]]

The classification matrix for decision tree is:

precision

recall f1-score support

Iris-setosa	1.00	1.00	1.00	19
Iris-versicolor	1.00	0.93	0.97	15
Iris-virginica	0.92	1.00	0.96	11
accuracy			0.98	45
macro avg	0.97	0.98	0.97	45
weighted avg	0.98	0.98	0.98	45

In [ ]: